

IPC

CPE 545

Purposes for IPC

- Data Transfer
- Sharing Data
- Event notification
- Resource Sharing and Synchronization
- Process Control


IPC Mechanisms

- *Mechanisms used for communication and synchronization*
 - Message Passing
 - message passing interfaces, mailboxes and message queues
 - sockets, STREAMS, pipes
 - *Shared Memory*. Non-message passing systems
 - *Synchronization* – primitives such as semaphores to higher level mechanisms such as monitors
 - *Event Notification* - UNIX *signals*

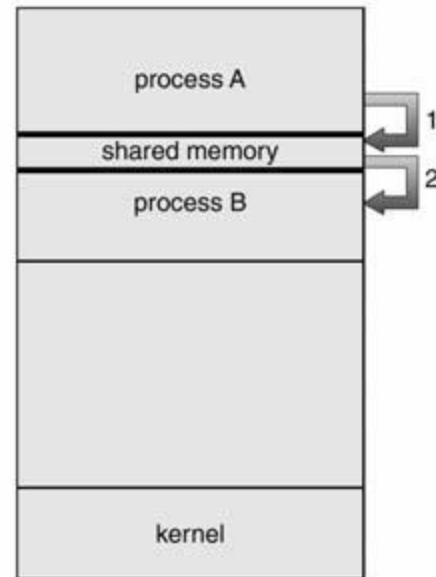
Message Passing

- In a *Message system* there are no shared variables. IPC facility provides two operations for fixed or variable sized message:
 - *send(message)*
 - *receive(message)*
- If processes P and Q wish to communicate, they need to:
 - establish a *communication link*
 - exchange messages via *send and receive*
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., syntax and semantics, abstractions)

Shared Memory



Message passing



Shared Memory

Mailboxes

- **Indirect communications :**

- Messages sent to and received from *mailboxes* (or *ports*)
 - mailboxes can be viewed as objects into which messages placed by processes and from which messages can be removed by other processes
- Each mailbox has a unique ID
- Two processes can communicate only if they have a shared mailbox
 - send** (A, message) : send a *message* to mailbox A
 - receive** (A, message) : receive a *message* from mailbox A
- A communications link is only established between a pair of processes if they have a shared mailbox
- A pair of processes can communicate via several different mailboxes if desired
- A link can be either unidirectional or bidirectional
- A link may be associated with more than two processes
 - allows *one-to-many*, *many-to-one*, *many-to-many* communications


Mailboxes- system ownership

- one-to-many : any of several processes may receive from the mailbox
 - which of the receivers gets the message?
 - only allow one process at a time to wait on a receive
- many-to-one : many processes sending to one receiving process
 - file server, network server, mail server etc.
 - receiver can identify the sender from the message header contents
- many-to-many :
 - e.g. multiple senders requesting service and a pool of receiving servers offering service - a *server farm*

Mailboxes- Process mailbox ownership

- only the process may receive messages from the mailbox
- other processes may send to the mailbox
- mailbox can be created with the process and destroyed when the process dies
 - process sending to a dead process's mailbox will need to be signaled
- *or* through separate **create_mailbox** and **destroy_mailbox** calls
 - possibly declare variables of type 'mailbox'

Mailboxes



Mailboxes

- A mailbox can correspond to a semaphore:
non-blocking send + *blocking* receive
 - equivalent to :
signal by sender (semGive) + *wait* by receiver (semTake)
- *Mutual Exclusion* :
 - initialize :
 create_mailbox (mutex)
 send (mutex, null-message)
 - for each process :
 while (TRUE)
 {
 receive (mutex, null-message);
 critical section
 send (mutex, null-message);
 }
 - mutual exclusion just depends on whether mailbox is empty or not
 - message is just a token, possession of which gives right to enter C.S.

Homework

Synchronization with Mutex

- Task A has higher priority than Task B
- In a special corner case condition in task A, if procedure `getAlarmStat()` is called, a call to `UpdateAlarms()` must be called by task B first, or the status of alarms are useless.
- Use Mutex or Mailbox to make sure the above order of execution is enforced

Producer / Consumer

- *Producer / Consumer* problem using messages :
 - *Binary* semaphores : one message token
 - *General (counting)* semaphores : more than one message token
 - message blocks used to buffer data items
 - scheme uses two mailboxes
 - *mayproduce* and *mayconsume*
- *producer* :
 - get a message block from *mayproduce*
 - put data item in block
 - send message to *mayconsume*
- *consumer* :
 - get a message from *mayconsume*
 - consume data in block
 - return empty message block to *mayproduce* mailbox

Producer / Consumer

- parent process creates message slots
 - buffering capacity depends on number of slots created
 - - slot = *empty message*
 - capacity = *buffering capacity*
 - create_mailbox (mayproduce);
 - create_mailbox (mayconsume);
 - for (i=0; i<capacity; i++) send (mayproduce, slot);
 - start producer and consumer processes*
- producer :
 - while (TRUE) {
 - receive (mayproduce, slot);
 - slot = *new data item*
 - send (mayconsume, slot);}
- consumer :
 - while (TRUE) {
 - receive (mayconsume, slot);
 - consume data item in slot*
 - send (mayproduce, slot);}

Event/Signals

- *Signals*

- The mechanism whereby processes are made aware of events occurring
- Asynchronous - can be received by a process at any time in its execution
- Examples of Linux signal types:
 - SIGINT : interrupt from keyboard
 - SIGFPE : floating point exception
 - SIGKILL : terminate receiving process
 - SIGCHLD : child process stopped or terminated
 - SIGSEGV : segment access violation
- Default action is usually for kernel to terminate the receiving process

Event/Signals

- Process can request some other action
 - ignore the signal - process will not know it happened
 - SIGKILL and SIGSTOP cannot be ignored
 - restore signal's default action
 - execute a pre-arranged signal-handling function
 - process can register a function to be called
 - like an interrupt service routine
 - when the handler returns, control is passed back to the main process code and normal execution continues
- to set up a signal handler:
`void (*signal(int signum, void (*handler)(int)))(int);`
- signal is a call which takes two parameters
 - signum : the signal number
 - handler : a pointer to a function which takes a single integer parameter and returns nothing (void)
- return value is itself a pointer to a function which:
 - takes a single integer parameter and returns nothing

Signals: Changing the behavior of ^C

```
main () {  
    int c;  
    old_handler = signal (SIGINT, ctrl_c );  
    while ((c = getchar()) != '\n'); //stuck here till hit <RET>  
    printf("ctrl_c count = %d\n", ctrl_c_count);  
    (void) signal (SIGINT, old_handler);  
  
    for (;;) ;  
}
```



```
void ctrl_c(int signum) {  
    (void) signal (SIGINT, ctrl_c);    // signals are  
    automatically reset  
    ++ctrl_c_count;  
}
```

- gets characters until a newline typed, then goes into an infinite loop
- uses signals to count ctrl-c's typed at keyboard until newline typed

Homework

How would you design an event handler framework like Linux Signals?